

## SAMPLE CLOUD COMPUTING QUESTIONS

---

1. Is it possible to have an application-level virtualization? Support your answer by explaining the reasons and examples.
2. Design a WordCount system using the MapReduce techniques. Draw a figure and explain how your design works based on the given two sentences.

Input:

This is an apple.

Apple is red in color.

3. What is serverless computing? Is it the same as function as a service? Why is it called serverless?
4. Considering using a control theory to optimize cloud resource management, what are the sources of instability when applying the control theory?

**Sample Multiple Choice Question**

- 
1. Which component in GFS is primarily responsible for managing metadata, such as namespace and file-to-chunk mapping?
- A. The chunk server
  - B. The client interface
  - C. The replication manager
  - D. The single master node
  - E. The distributed lock service
  - F. The metadata proxy
- 
2. Which statement best describes the role of chunk servers in GFS?
- A. They maintain the file namespace and metadata
  - B. They manage client authentication and access control
  - C. They store the actual data chunks and handle read/write requests
  - D. They provide encryption services for data in transit
  - E. They balance the overall network load
  - F. They directly communicate with the master for all operations
- 
3. Why does GFS use a single master node?
- A. To eliminate the need for distributed consensus
  - B. To centralize metadata management and simplify coordination
  - C. To increase data replication performance
  - D. To allow clients to access data directly without authentication
  - E. To balance network traffic between different data centers
  - F. To store both metadata and actual file contents
- 
4. What is the primary goal of the Ray framework?
- A. To provide a specialized deep learning framework for neural networks
  - B. To optimize traditional batch processing for big data analytics
  - C. To support distributed execution of AI and machine learning applications
  - D. To replace Kubernetes for container orchestration
  - E. To provide a real-time streaming system for IoT data processing
  - F. To offer a cloud-based service for GPU-based computations
- 
5. What component in Ray is responsible for managing distributed task execution?
- A. The Ray Client
  - B. The Ray Master Node
  - C. The Global Control Store (GCS)
  - D. The Ray Trainer Module
  - E. The Cluster State Manager
  - F. The AI Execution Engine

6. Which of the following is NOT a key feature of Ray?
- A. Support for heterogeneous computing environments
  - B. Built-in distributed reinforcement learning support
  - C. Automatic scheduling of machine learning training jobs
  - D. Direct execution of SQL queries on large datasets
  - E. Fault-tolerant task execution
  - F. Actor-based execution for dynamic workloads
- 
7. What is a Borgcell?
- A. A collection of machines managed by a single Borg master
  - B. A lightweight virtual machine used in Borg
  - C. A distributed file system for storing application data
  - D. A backup cluster used for fault tolerance
  - E. A specialized scheduling algorithm within Borg
  - F. A network partitioning mechanism
- 
8. What component of Borg is responsible for tracking job and task states?
- A. The Borgmaster
  - B. The Node Agent
  - C. The Borg UI
  - D. The Borg Database
  - E. The Priority Queue Manager
  - F. The Job Monitor